
Report for the Deep Learning Course Assignment 2

Georgios Methenitis
georgios.methenitis@cwil.nl

Abstract

In this assignment I got familiar with TensorFlow. I created a single hidden layer network with ReLU activation function to learn features for the CIFAR10 dataset.

1 Task 1

- 1.
2. There are two ways to declare variable in TensorFlow:
 - `var = tf.get_variable("var_name", shape=[..., ...])`
 - `var = tf.Variable(tf.zeros([..., ...]), name="var_name")`
3. The command `tf.shape(x)` returns the shape of the input of tensor `x` in an 1-D vector, while `x.get_shape()` return the size of the tensor.
- 4.
- 5.
- 6.
- 7.
- 8.
- 9.
- 10.

2 Task 2

In this task I implemented the following functions in `mlp.py`, more specifically the inference, loss, and accuracy.

```
def inference(self, x):
    with tf.variable_scope('hidden', reuse=None):
        W = tf.get_variable("weights", shape=[x.get_shape()[1], self.n_hidden[0]],
                            initializer=self.weight_initializer, regularizer=self.weight_regularizer)
        tf.histogram_summary("hidden_weights", W)
        b = tf.Variable(tf.zeros([self.n_hidden[0]]), name="bias")
        tf.histogram_summary("hidden_bias", b)

    with tf.variable_scope('output', reuse=None):
        w_out = tf.get_variable("weights", shape=[self.n_hidden[0], self.n_classes],
                                initializer=self.weight_initializer, regularizer=self.weight_regularizer)
        tf.histogram_summary("output_weights", w_out)
        b_out = tf.Variable(tf.zeros([self.n_classes]), name="bias")
        tf.histogram_summary("output_bias", b_out)

    with tf.name_scope('relu_layer'):
```

```

        input = self.activation_fn(tf.matmul(x, W) + b)
    with tf.name_scope('linear_layer'):
        output = tf.matmul(input, w_out) + b_out

    return output

def loss(self, logits, labels):

    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits, labels, name=None))
    with tf.variable_scope("hidden", reuse=True):
        loss += 0.5 * tf.reduce_sum(tf.get_variable("weights")**2)
    with tf.variable_scope("output", reuse=True):
        loss += 0.5 * tf.reduce_sum(tf.get_variable("weights")**2)

    return loss

def accuracy(self, logits, labels):
    correct_prediction = tf.equal(tf.argmax(labels, 1), tf.argmax(logits, 1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    return accuracy

```

I also implemented the function train in the file train_mlp.py:

```

def train():
    # placeholder for input and output tensors allowing to create operations and the computation graph
    x = tf.placeholder(tf.float32, shape=[None, 3072])
    y_ = tf.placeholder(tf.float32, shape=[None, 10])

    # model, namescopes for summary
    model = MLP()
    x_ = model.inference(x)
    with tf.name_scope('loss'):
        loss = model.loss(x_, y_)
    with tf.name_scope('accuracy'):
        accuracy = model.accuracy(x_, y_)
    with tf.name_scope('train_step'):
        train_step = tf.train.GradientDescentOptimizer(LEARNING_RATE_DEFAULT).minimize(loss)

    # session
    init = tf.initialize_all_variables()
    sess = tf.Session()
    sess.run(init)

    # merged summary for the graph, the loss and accuracy
    tf.scalar_summary('accuracy', accuracy)
    tf.scalar_summary('loss', loss)
    tf.histogram_summary("logits", x_) # logits histogram
    merged = tf.merge_all_summaries()

    train_writer = tf.train.SummaryWriter(LOG_DIR_DEFAULT + '/train', sess.graph)
    test_writer = tf.train.SummaryWriter(LOG_DIR_DEFAULT + '/test', sess.graph)

    for i in range(1, MAX_STEPS_DEFAULT+1):
        batch_xs, batch_ys = cifar10.train.next_batch(BATCH_SIZE_DEFAULT)
        __, Summary, l, acc = sess.run([train_step, merged, loss, accuracy],
            feed_dict={x: batch_xs, y_: batch_ys})
        train_writer.add_summary(Summary, i)

        if i % 100 == 0.0:
            batch_xs, batch_ys = cifar10.test.images, cifar10.test.labels
            __, Summary, l, acc = sess.run([train_step, merged, loss, accuracy],
                feed_dict={x: batch_xs, y_: batch_ys})
            test_writer.add_summary(Summary, i)

```




Figure 1: Recorded accuracy and loss for the training set.



Figure 2: Recorded accuracy and loss for the test set.

```
train_writer.close()
```

Figure 1 presents the accuracy and the loss in the training set recorded for every epoch.

Figure 2 presents the accuracy and the loss in the test set recorded for every 100-epochs.

Figure 3 presents the graph of the created model as this is illustrated by TensorBoard.

Figure 4 presents the histograms of the weights and the biases in every unit of the model, in the hidden and the output units respectively.

Figure 5 presents the histogram of the logits.

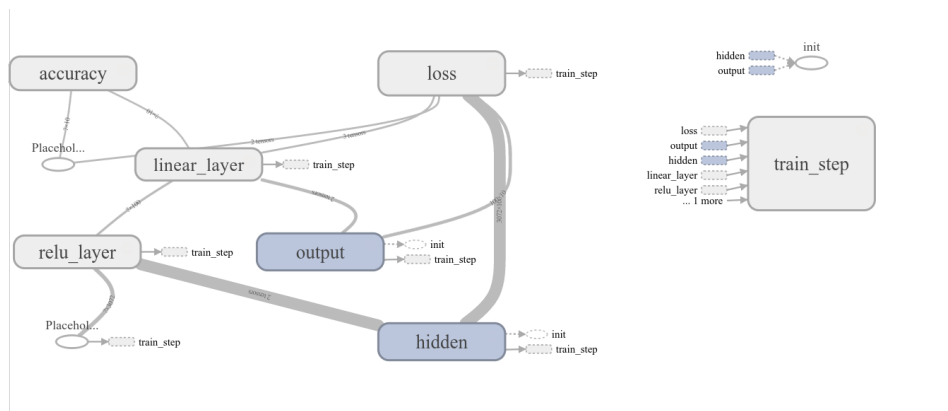


Figure 3: Graph of the model.



Figure 4: Histograms of weights and bias in the hidden layer, and weights and bias in the output linear layer.




Figure 5: Histograms of logits.

3 Conclusion

Should contain conclusion of this study. For example, you can try to answer the following questions. What was done during this assignment? What features of TensorBoard were positive and what were negative for implementing MLP model and performing the experiments? What are the main insights you got from the study of the MLP model on CIFAR10 dataset?