
Report for the Deep Learning Course Assignment 2

Georgios Methenitis
georgios.methenitis@cwi.nl

Abstract

In this assignment I got familiar with TensorFlow. I created a single hidden layer network with ReLU activation function to learn features for the CIFAR10 dataset. The accuracy obtained in the test set, 0.46, is similar to the neural network model of the previous assignment.

1 Task 1

1. TensorFlow *variables* are stored structures containing tensors, *constants* are constant variables in Tensorflow, *placeholders* are used in the place of variables at graph initialization when there are no variables. The differences are that constants cannot change their values in contrast with variables, and that placeholder are replacing variables at initialization phase.
2. There are two ways to declare variable in TensorFlow:
 - `var = tf.get_variable("var_name", shape=[..., ...])`
 - `var = tf.Variable(tf.zeros([..., ...]), name="var_name")`
3. The command `tf.shape(x)` returns the shape of the input of tensor `x` in an 1-D vector, while `x.get_shape()` return the size of the tensor.
- 4.
- 5.
6. A computational *graph*, *tensors* for the variables and *operator*.
7. In variable scopes, variables can be created, adopting the prefix of the scope in their name, variable scopes enable variable sharing in different places in the code. The differences between name scopes and variable scopes is in the former creates an hierarchy for operations and variables in the graph, while the latter helps mostly in variable sharing.
8. You can freeze a variable tensor during training by setting `trainable=True` upon creating the variable, this way it will not get updated during optimization.
9. Yes Tensorflow uses automatic differentiation, while this can be a disadvantage for higher numbers of hidden layer due to dependencies. The advantages are obvious, there is no need to compute difficult derivatives of high complexity.
10. One way is like the assignment, reading using a python script the data for every train step and feeding the data in the graph. The other way is to read all the data from csv files at the beginning.

2 Task 2

In this task I implemented the following functions in `mlp.py`, more specifically the `inference`, `loss`, and `accuracy`.

```
def inference(self, x):  
    with tf.variable_scope('hidden', reuse=None):
```

```

        W = tf.get_variable("weights", shape=[x.get_shape()[1], self.n_hidden[0]],
            initializer=self.weight_initializer, regularizer=self.weight_regularizer)
        tf.histogram_summary("hidden_weights", W)
        b = tf.Variable(tf.zeros([self.n_hidden[0]]), name="bias")
        tf.histogram_summary("hidden_bias", b)

    with tf.variable_scope('output', reuse=None):
        w_out = tf.get_variable("weights", shape=[self.n_hidden[0], self.n_classes],
            initializer=self.weight_initializer, regularizer=self.weight_regularizer)
        tf.histogram_summary("output_weights", w_out)
        b_out = tf.Variable(tf.zeros([self.n_classes]), name="bias")
        tf.histogram_summary("output_bias", b_out)

    with tf.name_scope('relu_layer'):
        input = self.activation_fn(tf.matmul(x, W) + b)
    with tf.name_scope('linear_layer'):
        output = tf.matmul(input, w_out) + b_out

    return output

def loss(self, logits, labels):

    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits, labels, name=None))
    with tf.variable_scope("hidden", reuse=True):
        loss += 0.5 * tf.reduce_sum(tf.get_variable("weights")**2)
    with tf.variable_scope("output", reuse=True):
        loss += 0.5 * tf.reduce_sum(tf.get_variable("weights")**2)

    return loss

def accuracy(self, logits, labels):
    correct_prediction = tf.equal(tf.argmax(labels, 1), tf.argmax(logits, 1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    return accuracy

```

I also implemented the function train in the file train_mlp.py:

```

def train():
    # placeholder for input and output tensors allowing to create operations and the computation graph
    x = tf.placeholder(tf.float32, shape=[None, 3072])
    y_ = tf.placeholder(tf.float32, shape=[None, 10])

    # model, namespaces for summary
    model = MLP()
    x_ = model.inference(x)
    with tf.name_scope('loss'):
        loss = model.loss(x_, y_)
    with tf.name_scope('accuracy'):
        accuracy = model.accuracy(x_, y_)
    with tf.name_scope('train_step'):
        train_step = tf.train.GradientDescentOptimizer(LEARNING_RATE_DEFAULT).minimize(loss)

    # session
    init = tf.initialize_all_variables()
    sess = tf.Session()
    sess.run(init)

    # merged summary for the graph, the loss and accuracy
    tf.scalar_summary('accuracy', accuracy)
    tf.scalar_summary('loss', loss)
    tf.histogram_summary("logits", x_) # logits histogram
    merged = tf.merge_all_summaries()

    train_writer = tf.train.SummaryWriter(LOG_DIR_DEFAULT + '/train', sess.graph)

```




Figure 1: Recorded accuracy and loss for the training set.



Figure 2: Recorded accuracy and loss for the test set.

```
test_writer = tf.train.SummaryWriter(LOG_DIR_DEFAULT + '/test', sess.graph)

for i in range(1, MAX_STEPS_DEFAULT+1):
    batch_xs, batch_ys = cifar10.train.next_batch(BATCH_SIZE_DEFAULT)
    __, Summary, l, acc = sess.run([train_step, merged, loss, accuracy],
    feed_dict={x: batch_xs, y_: batch_ys})
    train_writer.add_summary(Summary, i)

    if i % 100 == 0.0:
        batch_xs, batch_ys = cifar10.test.images, cifar10.test.labels
        __, Summary, l, acc = sess.run([train_step, merged, loss, accuracy],
        feed_dict={x: batch_xs, y_: batch_ys})
        test_writer.add_summary(Summary, i)

train_writer.close()
```

Figure 1 presents the accuracy and the loss in the training set recorded for every epoch. The accuracy is higher than 0.45 similar to the performance in the training set of the previous assignment, while the loss converges close to zero after 500 epochs. Figure 2 presents the accuracy and the loss in the test set recorded every 100-epochs. As you can see the accuracy drops after 1200 epochs, the model is overfitting.

Figure 3 presents the graph of the created model as this is illustrated by TensorBoard. Using different namespaces where necessary we can illustrate the structure of the model. You can see the two placeholder, one given as input to the relu layer and the second as the output of the linear layer, also connected to the accuracy module. The two layer are connected to the loss module that then updates all weights in the output and hidden layers.

Figure 4 presents the histograms of the weights and the biases in every unit of the model, in the hidden and the output units respectively.

Figure 5 presents the histogram of the logits. We can observe that only the variance of the logits change.

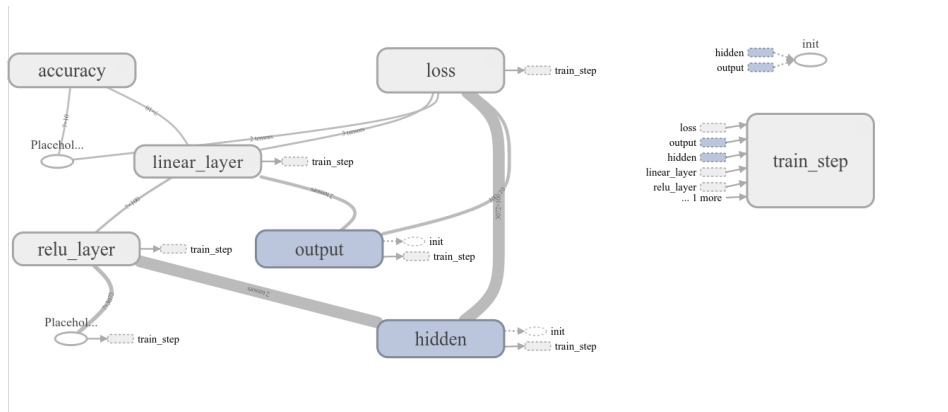


Figure 3: Graph of the model.



Figure 4: Histograms of weights and bias in the hidden layer, and weights and bias in the output linear layer.




Figure 5: Histograms of logits.

3 Conclusion

Although I did not have time to finish all the questions in the assignment I found it very useful. Tensorflow provides an easy to use library to create and train models and visualize everything is happening during the training of the model.